


Figure 2.2 — Time-Aware Transformer Architecture



Input (48×74) projected to $d = 64$ via a shared linear layer, summed with learnable sinusoidal TAPE, prepended with a [CLS] token, and passed through 4 pre-norm encoder layers (total 205 153 parameters). The contextualised CLS vector feeds a two-layer MLP ($64 \rightarrow 32 \rightarrow 1$) that emits the mortality logit.